

Starbucks_Capstone_notebook

July 12, 2021

1 Starbucks Capstone Challenge

1.0.1 Introduction

This data set contains simulated data that mimics customer behavior on the Starbucks rewards mobile app. Once every few days, Starbucks sends out an offer to users of the mobile app. An offer can be merely an advertisement for a drink or an actual offer such as a discount or BOGO (buy one get one free). Some users might not receive any offer during certain weeks.

Not all users receive the same offer, and that is the challenge to solve with this data set.

Your task is to combine transaction, demographic and offer data to determine which demographic groups respond best to which offer type. This data set is a simplified version of the real Starbucks app because the underlying simulator only has one product whereas Starbucks actually sells dozens of products.

Every offer has a validity period before the offer expires. As an example, a BOGO offer might be valid for only 5 days. You'll see in the data set that informational offers have a validity period even though these ads are merely providing information about a product; for example, if an informational offer has 7 days of validity, you can assume the customer is feeling the influence of the offer for 7 days after receiving the advertisement.

You'll be given transactional data showing user purchases made on the app including the timestamp of purchase and the amount of money spent on a purchase. This transactional data also has a record for each offer that a user receives as well as a record for when a user actually views the offer. There are also records for when a user completes an offer.

Keep in mind as well that someone using the app might make a purchase through the app without having received an offer or seen an offer.

1.0.2 Example

To give an example, a user could receive a discount offer buy 10 dollars get 2 off on Monday. The offer is valid for 10 days from receipt. If the customer accumulates at least 10 dollars in purchases during the validity period, the customer completes the offer.

However, there are a few things to watch out for in this data set. Customers do not opt into the offers that they receive; in other words, a user can receive an offer, never actually view the offer, and still complete the offer. For example, a user might receive the "buy 10 dollars get 2 dollars off offer", but the user never opens the offer during the 10 day validity period. The customer spends 15 dollars during those ten days. There will be an offer completion record in the data set; however, the customer was not influenced by the offer because the customer never viewed the offer.

1.0.3 Cleaning

This makes data cleaning especially important and tricky.

You'll also want to take into account that some demographic groups will make purchases even if they don't receive an offer. From a business perspective, if a customer is going to make a 10 dollar purchase without an offer anyway, you wouldn't want to send a buy 10 dollars get 2 dollars off offer. You'll want to try to assess what a certain demographic group will buy when not receiving any offers.

1.0.4 Final Advice

Because this is a capstone project, you are free to analyze the data any way you see fit. For example, you could build a machine learning model that predicts how much someone will spend based on demographics and offer type. Or you could build a model that predicts whether or not someone will respond to an offer. Or, you don't need to build a machine learning model at all. You could develop a set of heuristics that determine what offer you should send to each customer (i.e., 75 percent of women customers who were 35 years old responded to offer A vs 40 percent from the same demographic to offer B, so send offer A).

2 Data Sets

The data is contained in three files:

- `portfolio.json` - containing offer ids and meta data about each offer (duration, type, etc.)
- `profile.json` - demographic data for each customer
- `transcript.json` - records for transactions, offers received, offers viewed, and offers completed

Here is the schema and explanation of each variable in the files:

portfolio.json * `id` (string) - offer id * `offer_type` (string) - type of offer ie BOGO, discount, informational * `difficulty` (int) - minimum required spend to complete an offer * `reward` (int) - reward given for completing an offer * `duration` (int) - time for offer to be open, in days * `channels` (list of strings)

profile.json * `age` (int) - age of the customer * `became_member_on` (int) - date when customer created an app account * `gender` (str) - gender of the customer (note some entries contain 'O' for other rather than M or F) * `id` (str) - customer id * `income` (float) - customer's income

transcript.json * `event` (str) - record description (ie transaction, offer received, offer viewed, etc.) * `person` (str) - customer id * `time` (int) - time in hours since start of test. The data begins at time `t=0` * `value` - (dict of strings) - either an offer id or transaction amount depending on the record

Note: If you are using the workspace, you will need to go to the terminal and run the command `conda update pandas` before reading in the files. This is because the version of pandas in the workspace cannot read in the `transcript.json` file correctly, but the newest version of pandas can. You can access the terminal from the orange icon in the top left of this notebook.

You can see how to access the terminal and how the install works using the two images below. First you need to access the terminal:

Then you will want to run the above command:

Finally, when you enter back into the notebook (use the jupyter icon again), you should be able to run the below cell without any errors.

```

In [3]: #regular libraries
import pandas as pd
import numpy as np
import math
import json

#plot libraries
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import seaborn as sns
% matplotlib inline

#machine learning library
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
#warnings ignore
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

# read in the json files
portfolio = pd.read_json('data/portfolio.json', orient='records', lines=True)
profile = pd.read_json('data/profile.json', orient='records', lines=True)
transcript = pd.read_json('data/transcript.json', orient='records', lines=True)
pd.set_option('display.max_columns', None)

In [4]: print('\n'.join(f'{m.__name__}=={m.__version__}' for m in globals().values() if getattr(
pandas==0.23.3
numpy==1.12.1
json==2.0.9
seaborn==0.8.1

```

3 Data Understanding

```
In [5]: portfolio.head()
```

```
Out[5]:
```

	channels	difficulty	duration	\
0	[email, mobile, social]	10	7	
1	[web, email, mobile, social]	10	5	
2	[web, email, mobile]	0	4	
3	[web, email, mobile]	5	7	

	4	[web, email]	20	10	
			id	offer_type	reward
0	ae264e3637204a6fb9bb56bc8210ddfd			bogo	10
1	4d5c57ea9a6940dd891ad53e9dbe8da0			bogo	10
2	3f207df678b143eea3cee63160fa8bed			informational	0
3	9b98b8c7a33c4b65b9aebfe6a799e6d9			bogo	5
4	0b1e1539f2cc45b7b9fa7c272da2e1d7			discount	5

```
In [6]: portfolio.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 6 columns):
channels      10 non-null object
difficulty    10 non-null int64
duration      10 non-null int64
id            10 non-null object
offer_type    10 non-null object
reward        10 non-null int64
dtypes: int64(3), object(3)
memory usage: 560.0+ bytes
```

```
In [7]: profile.head()
```

```
Out[7]:
```

	age	became_member_on	gender	id	income
0	118	20170212	None	68be06ca386d4c31939f3a4f0e3dd783	NaN
1	55	20170715	F	0610b486422d4921ae7d2bf64640c50b	112000.0
2	118	20180712	None	38fe809add3b4fcf9315a9694bb96ff5	NaN
3	75	20170509	F	78afa995795e4d85b5d9ceeca43f5fef	100000.0
4	118	20170804	None	a03223e636434f42ac4c3df47e8bac43	NaN

```
In [8]: profile.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17000 entries, 0 to 16999
Data columns (total 5 columns):
age      17000 non-null int64
became_member_on  17000 non-null int64
gender    14825 non-null object
id        17000 non-null object
income    14825 non-null float64
dtypes: float64(1), int64(2), object(2)
memory usage: 664.1+ KB
```

```
In [9]: profile.describe(include='all')
```

```

Out[9]:
      age  became_member_on  gender  \
count  17000.000000      1.700000e+04  14825
unique      NaN              NaN      3
top         NaN              NaN      M
freq        NaN              NaN    8484
mean      62.531412      2.016703e+07   NaN
std       26.738580      1.167750e+04   NaN
min       18.000000      2.013073e+07   NaN
25%       45.000000      2.016053e+07   NaN
50%       58.000000      2.017080e+07   NaN
75%       73.000000      2.017123e+07   NaN
max      118.000000      2.018073e+07   NaN

      id      income
count      17000    14825.000000
unique      17000         NaN
top  0acca8aae113433999f7de6a5c32497c         NaN
freq         1         NaN
mean         NaN    65404.991568
std         NaN    21598.299410
min         NaN    30000.000000
25%         NaN    49000.000000
50%         NaN    64000.000000
75%         NaN    80000.000000
max         NaN   120000.000000

```

```
In [10]: transcript.head()
```

```

Out[10]:
      event      person  time  \
0  offer received  78afa995795e4d85b5d9ceeca43f5fef      0
1  offer received  a03223e636434f42ac4c3df47e8bac43      0
2  offer received  e2127556f4f64592b11af22de27a7932      0
3  offer received  8ec6ce2a7e7949b1bf142def7d0e0586      0
4  offer received  68617ca6246f4fbc85e91a2a49552598      0

      value
0  {'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}
1  {'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}
2  {'offer id': '2906b810c7d4411798c6938adc9daaa5'}
3  {'offer id': 'fafdc668e3743c1bb461111dcafc2a4'}
4  {'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}

```

```
In [11]: transcript.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 306534 entries, 0 to 306533
Data columns (total 4 columns):
event      306534 non-null object
person     306534 non-null object


```

```
time      306534 non-null int64
value     306534 non-null object
dtypes: int64(1), object(3)
memory usage: 9.4+ MB
```

4 Data Exploration

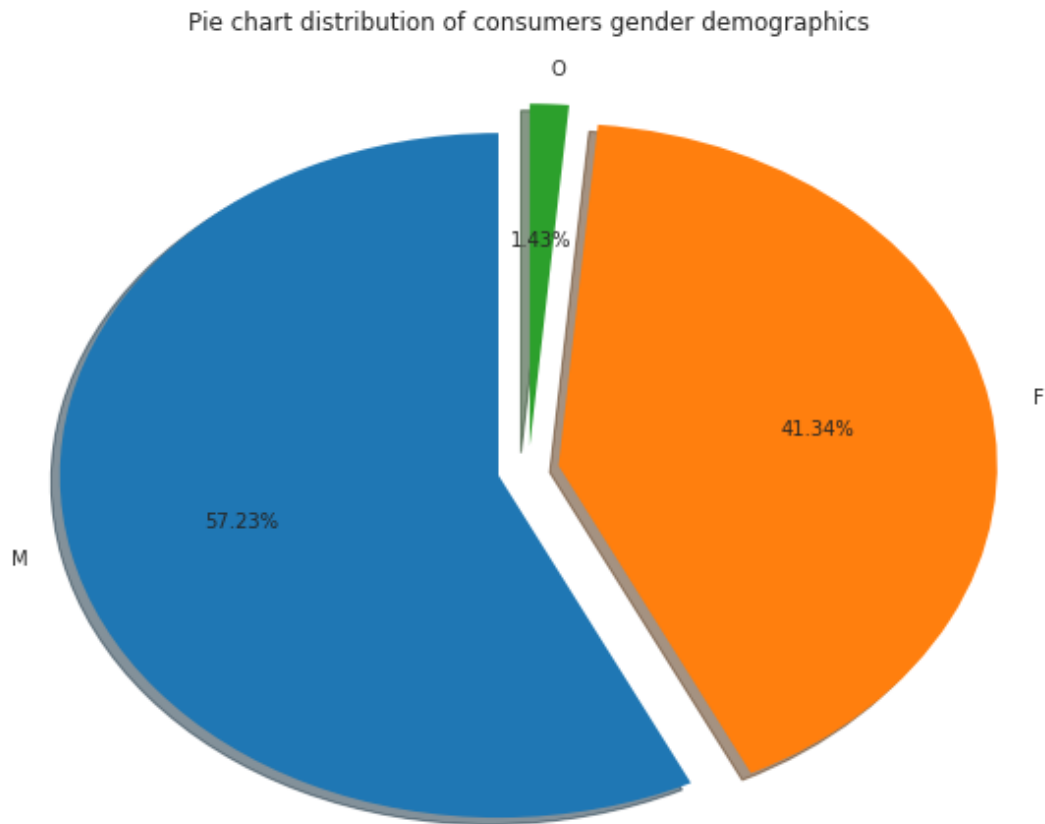
Below plot shows the distribution of consumer's age

```
In [12]: sns.set_style("whitegrid")
fig, ax = plt.subplots(figsize=(20,8))
sns.distplot(a=profile['age'], color='g', kde = True);
ax.set_title('Distribution of consumers age');
```



We can see abnormal amount of values near 120. This might be due to Null values replaced with this value. It will be delt in near future code.

```
In [13]: #pie plot of gender distribution
fig, ax = plt.subplots(figsize=(10,8))
plt.pie(profile['gender'].value_counts().values, labels= profile['gender'].value_counts(
ax.set_title('Pie chart distribution of consumers gender demographics');
```



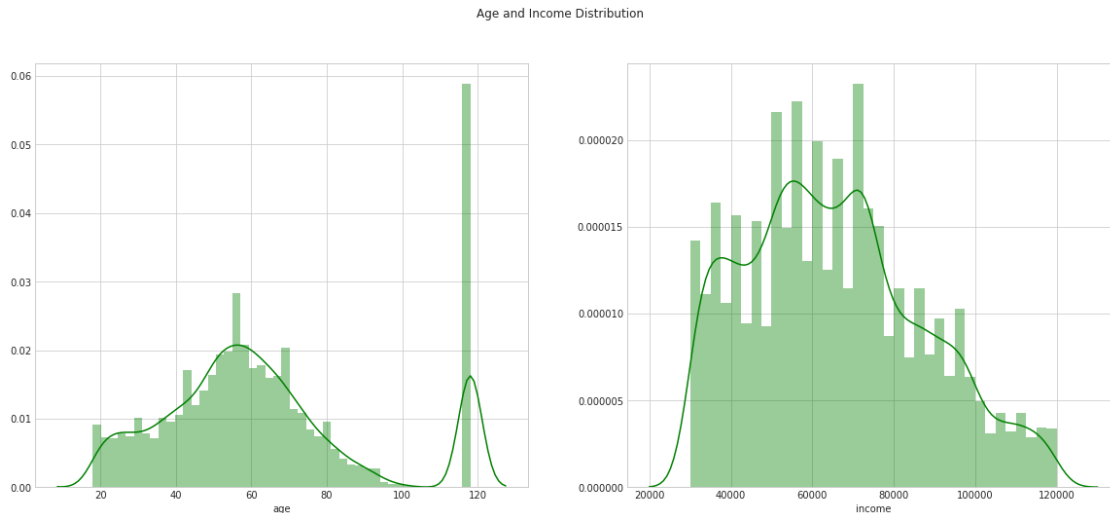
```
In [14]: profile[profile.isna().any(axis=1)].head()
```

```
Out[14]:
```

	age	became_member_on	gender	id	income
0	118	20170212	None	68be06ca386d4c31939f3a4f0e3dd783	NaN
2	118	20180712	None	38fe809add3b4fcf9315a9694bb96ff5	NaN
4	118	20170804	None	a03223e636434f42ac4c3df47e8bac43	NaN
6	118	20170925	None	8ec6ce2a7e7949b1bf142def7d0e0586	NaN
7	118	20171002	None	68617ca6246f4fbc85e91a2a49552598	NaN

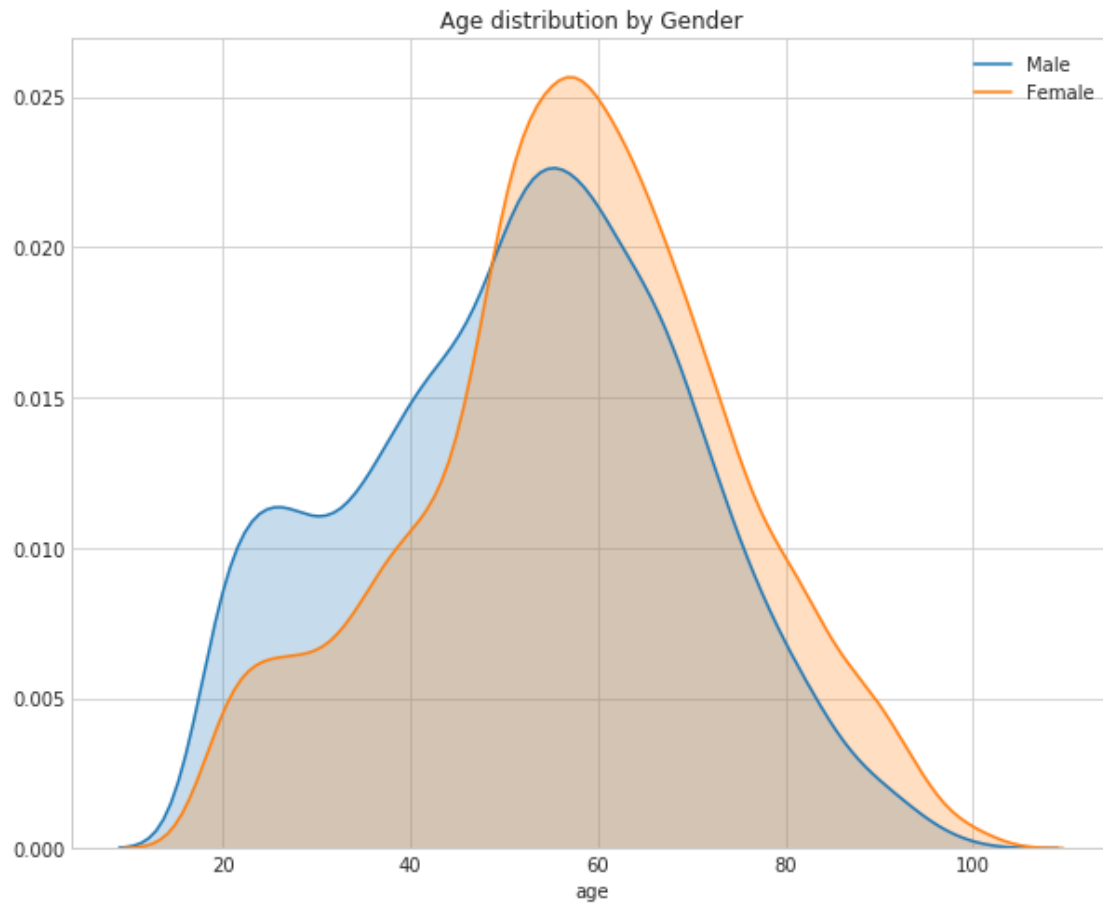
From above table we can confirm that Null values in age column is replaced with 118. Along with it None for gender column and NaN for income column

```
In [15]: #plotting age and income distribution
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(20, 8))
sns.distplot(a=profile['age'][profile['age'].notnull()], color='g', ax=axes[0]);
sns.distplot(a=profile['income'][profile['income'].notnull()], color='g', ax=axes[1]);
fig.suptitle('Age and Income Distribution');
```

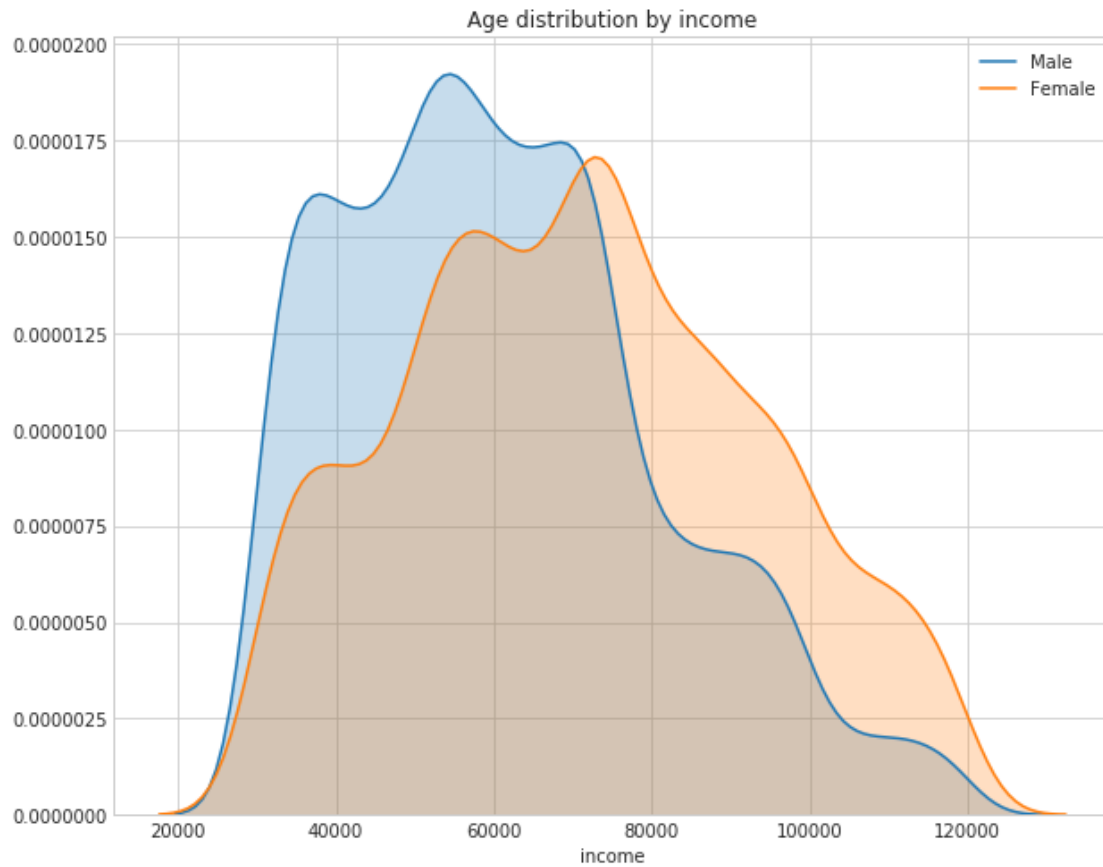


Above plots are made without Null values includes in it. From distribution it is clear: - Since the distribution of age is symmetric the null values can be replaced by mean - Since the distribution of income is left skewed the null values can be replaced by median It will be delt in near future code.

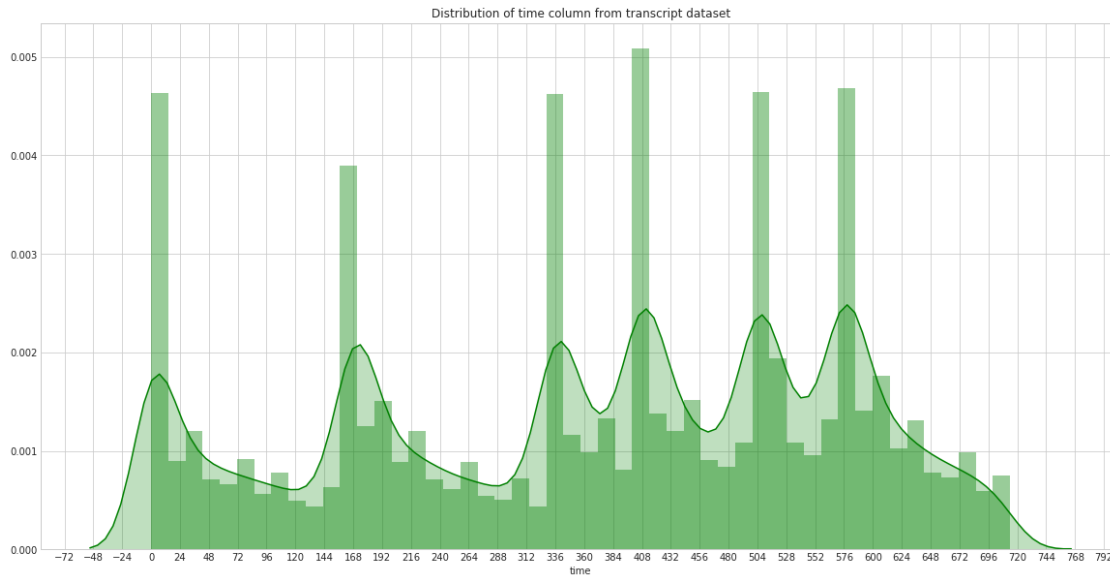
```
In [16]: #plotting age distribution by gender
fig, ax = plt.subplots(figsize=(10,8));
sns.distplot(profile[profile['gender']=='M']['age'],hist=False, kde_kws={"shade": True})
sns.distplot(profile[profile['gender']=='F']['age'],hist=False, kde_kws={"shade": True})
plt.title('Age distribution by Gender');
```

```
In [17]: #plotting age distribution by income
fig, ax = plt.subplots(figsize=(10,8));
sns.distplot(profile[profile['gender']=='M']['income'],hist=False, kde_kws={"shade": True});
sns.distplot(profile[profile['gender']=='F']['income'],hist=False,kde_kws={"shade": True});
plt.title('Age distribution by income');
plt.legend();
```



```
In [18]: #plot distribution of time column
fig, ax = plt.subplots(figsize=(20,10));
sns.distplot(transcript['time'], color='g', hist = True, kde_kws={"shade": True});
plt.title('Distribution of time column from transcript dataset');
ax.xaxis.set_major_locator(ticker.MultipleLocator(24))
ax.xaxis.set_major_formatter(ticker.ScalarFormatter())
```

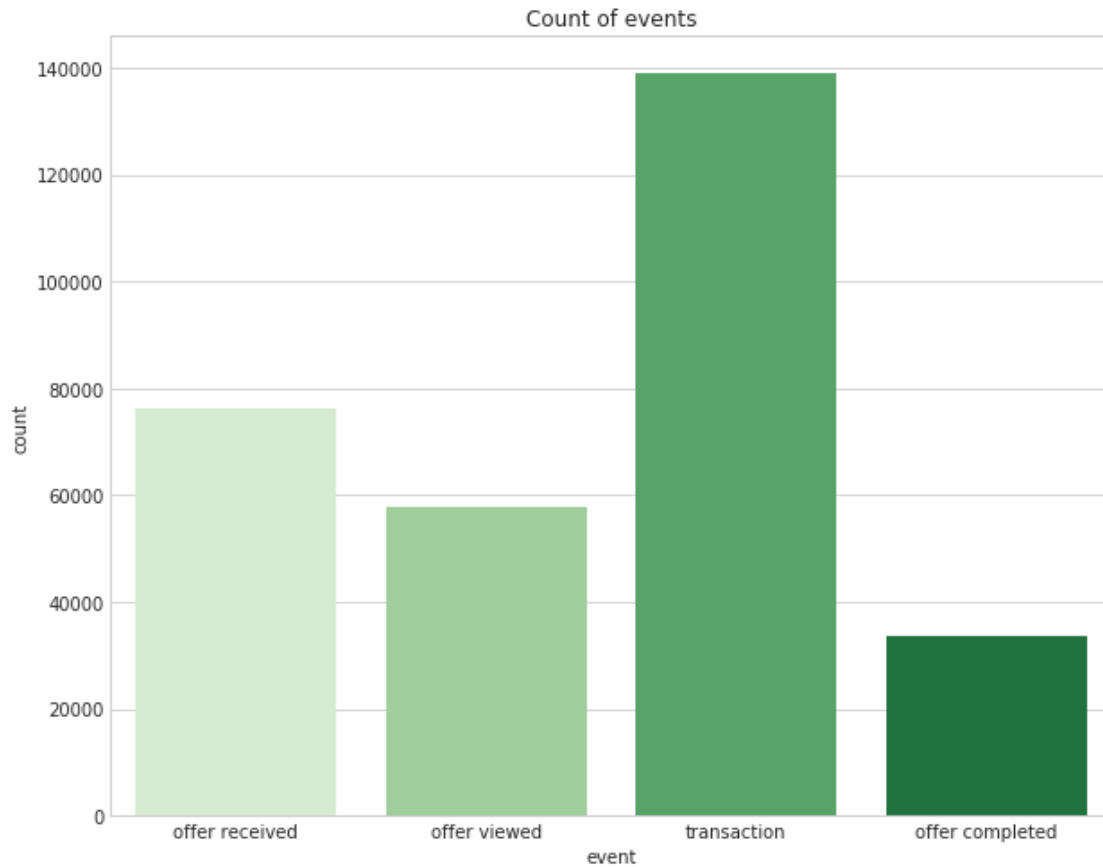


It is clear that distribution follows some pattern. From general knowledge, we know mornings will have high purchase volume from the stores as coffee is used boost the start of workday. Also it is also intuitive that Mondays will have high purchase volume then rest of the week.

Since no information is given about what and when the time = 0 starts from the transcript dataset. And since offer notifications, offer views drive purchases. **I will assume that time = 0 starts in the morning of Monday**

My assumptions can also be supported from the employees of starbuck's discussion: [Link](#)

```
In [19]: #plot count of events
fig, ax = plt.subplots(figsize=(10,8));
sns.countplot(transcript['event'],palette="Greens");
plt.title('Count of events');
```



5 Preparing Data

Glancing at portfolio dataset..

In [20]: `portfolio.head()`

```
Out[20]:
```

	channels	difficulty	duration	\
0	[email, mobile, social]	10	7	
1	[web, email, mobile, social]	10	5	
2	[web, email, mobile]	0	4	
3	[web, email, mobile]	5	7	
4	[web, email]	20	10	

	id	offer_type	reward
0	ae264e3637204a6fb9bb56bc8210ddfd	bogo	10
1	4d5c57ea9a6940dd891ad53e9dbe8da0	bogo	10
2	3f207df678b143eea3cee63160fa8bed	informational	0
3	9b98b8c7a33c4b65b9aebfe6a799e6d9	bogo	5
4	0b1e1539f2cc45b7b9fa7c272da2e1d7	discount	5

We can see that channels is in list format. It cannot be used directly, so some preprocessing is required. Below code creates a new column and enters 1 if a particular channel is present or else 0

```
In [21]: portfolio['channel_email'] = portfolio['channels'].apply(lambda x: 1 if 'email' in x else 0)
portfolio['channel_mobile'] = portfolio['channels'].apply(lambda x: 1 if 'mobile' in x else 0)
portfolio['channel_social'] = portfolio['channels'].apply(lambda x: 1 if 'social' in x else 0)
portfolio['channel_web'] = portfolio['channels'].apply(lambda x: 1 if 'web' in x else 0)
portfolio.drop(columns=['channels'], inplace=True)
```

```
In [22]: portfolio
```

```
Out[22]:
```

	difficulty	duration	id	offer_type	\
0	10	7	ae264e3637204a6fb9bb56bc8210ddfd	bogo	
1	10	5	4d5c57ea9a6940dd891ad53e9dbe8da0	bogo	
2	0	4	3f207df678b143eea3cee63160fa8bed	informational	
3	5	7	9b98b8c7a33c4b65b9aebfe6a799e6d9	bogo	
4	20	10	0b1e1539f2cc45b7b9fa7c272da2e1d7	discount	
5	7	7	2298d6c36e964ae4a3e7e9706d1fb8c2	discount	
6	10	10	fafdc668e3743c1bb461111dcafc2a4	discount	
7	0	3	5a8bc65990b245e5a138643cd4eb9837	informational	
8	5	5	f19421c1d4aa40978ebb69ca19b0e20d	bogo	
9	10	7	2906b810c7d4411798c6938adc9daaa5	discount	

	reward	channel_email	channel_mobile	channel_social	channel_web
0	10	1	1	1	0
1	10	1	1	1	1
2	0	1	1	0	1
3	5	1	1	0	1
4	5	1	0	0	1
5	3	1	1	1	1
6	2	1	1	1	1
7	0	1	1	1	0
8	5	1	1	1	1
9	2	1	1	0	1

As per general knowledge, We know that we try to maximize reward given the difficulty. So I feel the ratio between reward and difficulty feature will help model in prediction. This process is called **Feature Engineering**

```
In [23]: #create a column with reward/difficulty ratio and handle NaN values
portfolio['rew_by_diff'] = portfolio['reward']/portfolio['difficulty']
portfolio['rew_by_diff'].fillna(0, inplace=True)
```

As offer_type is categorical variable for future model prediction it is being encoded

```
In [24]: #one-hot encoding
offers = pd.get_dummies(portfolio['offer_type'], prefix='offer_type', prefix_sep='_')
```

```
In [25]: #concatenate with original dataset
portfolio = pd.concat([portfolio, offers], axis=1)
```

```
In [26]: #remove original offer_type column
portfolio.drop(columns=['offer_type'], inplace=True)
```

As id from portfolio can be confused with id from profile it is being renamed

```
In [27]: portfolio.rename(columns={'id': 'offer_id'}, inplace=True)
```

Glancing at profile dataset..

```
In [28]: profile.head()
```

```
Out[28]:
```

	age	became_member_on	gender	id	income
0	118	20170212	None	68be06ca386d4c31939f3a4f0e3dd783	NaN
1	55	20170715	F	0610b486422d4921ae7d2bf64640c50b	112000.0
2	118	20180712	None	38fe809add3b4fcf9315a9694bb96ff5	NaN
3	75	20170509	F	78afa995795e4d85b5d9ceeca43f5fef	100000.0
4	118	20170804	None	a03223e636434f42ac4c3df47e8bac43	NaN

From earlier comment => age column 118 will be converted to nan and later it will filled using relevant imputation technique like mean in this case

```
In [29]: #convert 118 to np.nan
profile['age'] = profile['age'].apply(lambda x: np.nan if x == 118 else x)
#convert dtype of became_member_on column
profile['became_member_on'] = pd.to_datetime(profile['became_member_on'], format='%Y%m%d')
```

As became_member_on is in pandas datetime dtype but it cannot be used in analysis or in prediction. So it engineered so as to convert it into an numerical. which will still support same information.

```
In [30]: profile['became_member_on'] = (profile['became_member_on'] - profile['became_member_on']
```

As gender is categorical variable for future model prediction it is being encoded

```
In [31]: #one-hot encoding
gender = pd.get_dummies(profile['gender'], prefix = 'gender', prefix_sep='_')
```

```
In [32]: #concatenate with orginal dataset. Delete gender column
profile = pd.concat([profile, gender], axis=1)
profile.drop(columns=['gender'], inplace=True)
```

```
In [33]: #For clarity and merging which will be used in later part, The id column is renamed
profile.rename(columns={'id': 'consumer_id'}, inplace=True)
```

```
In [34]: #From above discussion we impute age column with mean
profile['age'] = profile['age'].fillna(profile['age'].mean())
```

```
In [35]: #From above discussion we impute age column with median
profile['income'] = profile['income'].fillna(profile['income'].median())
```

```
In [36]: #Function for categorization of age column for further analysis
```

```
def age(x):  
    if x <= 30:  
        return "Young_Adult"  
    elif (x>30 and x<=60):  
        return "Adult"  
    else:  
        return "Old"
```

```
In [37]: #Function for categorization of income column for further analysis
```

```
def income(x):  
    if x <= 50000:  
        return "Lower"  
    elif (x > 50000 and x <= 90000):  
        return "Middle"  
    else:  
        return "Upper"
```

```
In [38]: #apply the age function
```

```
profile['Age_group'] = profile['age'].apply(age)  
#one hot encoding  
age_column = pd.get_dummies(profile['Age_group'], prefix='Age_group', prefix_sep='_')  
#concatanate with original dataset  
profile = pd.concat([profile, age_column], axis=1)  
#drop the Age_group column  
profile.drop(columns=['Age_group'], inplace=True)
```

```
In [39]: #Apply the income function
```

```
profile['Income_group'] = profile['income'].apply(income)  
#one hot enoding  
income_column = pd.get_dummies(profile['Income_group'], prefix='Income_group', prefix_sep='_')  
#concatanate with orifinal dataset  
profile = pd.concat([profile, income_column], axis=1)  
#drop the Income_group column  
profile.drop(columns=['Income_group'], inplace=True)
```

```
In [40]: profile.head()
```

```
Out[40]:
```

	age	became_member_on	consumer_id	income	\
0	54.393524	0.709819	68be06ca386d4c31939f3a4f0e3dd783	64000.0	
1	55.000000	0.793747	0610b486422d4921ae7d2bf64640c50b	112000.0	
2	54.393524	0.992320	38fe809add3b4fcf9315a9694bb96ff5	64000.0	
3	75.000000	0.756994	78afa995795e4d85b5d9ceeca43f5fef	100000.0	
4	54.393524	0.804717	a03223e636434f42ac4c3df47e8bac43	64000.0	

	gender_F	gender_M	gender_O	Age_group_Adult	Age_group_Old	\
0	0	0	0	1	0	
1	1	0	0	1	0	
2	0	0	0	1	0	

3	1	0	0	0	1
4	0	0	0	1	0

	Age_group_Young_Adult	Income_group_Lower	Income_group_Middle	\
0	0	0		1
1	0	0		0
2	0	0		1
3	0	0		0
4	0	0		1

	Income_group_Upper
0	0
1	1
2	0
3	1
4	0

```
In [41]: profile.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17000 entries, 0 to 16999
Data columns (total 13 columns):
age                17000 non-null float64
became_member_on   17000 non-null float64
consumer_id        17000 non-null object
income             17000 non-null float64
gender_F           17000 non-null uint8
gender_M           17000 non-null uint8
gender_O           17000 non-null uint8
Age_group_Adult    17000 non-null uint8
Age_group_Old      17000 non-null uint8
Age_group_Young_Adult 17000 non-null uint8
Income_group_Lower 17000 non-null uint8
Income_group_Middle 17000 non-null uint8
Income_group_Upper 17000 non-null uint8
dtypes: float64(3), object(1), uint8(9)
memory usage: 680.7+ KB
```

Glancing at transcript dataset..

```
In [42]: transcript.head()
```

```
Out[42]:
```

	event	person	time	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0	
1	offer received	a03223e636434f42ac4c3df47e8bac43	0	
2	offer received	e2127556f4f64592b11af22de27a7932	0	
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0	
4	offer received	68617ca6246f4fbc85e91a2a49552598	0	


```

                                value
0  {'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}
1  {'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}
2  {'offer id': '2906b810c7d4411798c6938adc9daaa5'}
3  {'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}
4  {'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}

```

In [43]: *#Function for categorization of time column for further analysis*

```

def period(x):
    hour = x%24
    period_internal = int(hour/6)
    if period_internal < 1:
        return "Morning"
    elif (period_internal < 2) and (period_internal >= 1):
        return "Afternoon"
    elif (period_internal < 3) and (period_internal >= 2):
        return "Evening"
    else:
        return "Night"

```

In [44]: *#Function for categorization of time column for further analysis*

```

def day(x):
    day_internal = int(x/24)
    if day_internal % 7 == 0:
        return "Monday"
    elif day_internal % 7 == 1:
        return "Tuesday"
    elif day_internal % 7 == 2:
        return "Wednesday"
    elif day_internal % 7 == 3:
        return "Thrusday"
    elif day_internal % 7 == 4:
        return "Friday"
    elif day_internal % 7 == 5:
        return "Saturday"
    else:
        return "Sunday"

```

In [45]: *#Apply period function*

```

transcript['period'] = transcript['time'].apply(period)
#one hot encoding
period_column = pd.get_dummies(transcript['period'], prefix='period', prefix_sep='_')
#concatanate with original dataset
transcript = pd.concat([transcript, period_column], axis=1)
#Will be used in later analysis
df_analysis_6 = transcript.copy()

```

```
#drop period column
transcript.drop(columns=['period'], inplace=True)
```

```
In [46]: #Apply day function
transcript['day'] = transcript['time'].apply(day)
#one hot encoding
day_column = pd.get_dummies(transcript['day'], prefix='day', prefix_sep='_')
#concatenate with original dataset
transcript = pd.concat([transcript, day_column], axis=1)
#Will be used in later analysis
df_analysis_4 = transcript.copy()
#drop period column
transcript.drop(columns=['day'], inplace=True)
```

```
In [47]: transcript.head()
```

```
Out[47]:
```

	event	person	time	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0	
1	offer received	a03223e636434f42ac4c3df47e8bac43	0	
2	offer received	e2127556f4f64592b11af22de27a7932	0	
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0	
4	offer received	68617ca6246f4fbc85e91a2a49552598	0	

	value	period_Afternoon	\
0	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	0	
1	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0	
2	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0	
3	{'offer id': 'fafdc668e3743c1bb461111dcafc2a4'}	0	
4	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}	0	

	period_Evening	period_Morning	period_Night	day_Friday	day_Monday	\
0	0	1	0	0	1	
1	0	1	0	0	1	
2	0	1	0	0	1	
3	0	1	0	0	1	
4	0	1	0	0	1	

	day_Saturday	day_Sunday	day_Thursday	day_Tuesday	day_Wednesday
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

Value column from this dataset is important as it contains the offer id, reward, amount according to the event. But preprocessing this column is tricky

```
In [48]: #As time column is measured in hours this transformation helps in finding trends in week
transcript['time'] = transcript['time']/24
```

```

transcript['offer_id'] = transcript['value'].apply(lambda x: x['offer_id'] if 'offer_id'
num_vals = ['reward', 'amount']
for i in num_vals:
    transcript[i] = transcript['value'].apply(lambda x: x[i] if i in x else None)

transcript.drop('value', axis=1, inplace=True)

In [49]: #Average_frequency is engineered feature which will give us idea about how frequent ite
#groupby person , event and apply difference of mean on time column and unstack with su
average_frequency = transcript.groupby(by=['person', 'event'])['time'].apply(lambda x:

/opt/conda/lib/python3.6/site-packages/numpy/core/fromnumeric.py:2889: RuntimeWarning: Mean of e
out=out, **kwargs)
/opt/conda/lib/python3.6/site-packages/numpy/core/_methods.py:80: RuntimeWarning: invalid value
ret = ret.dtype.type(ret / rcount)

In [50]: #merge it into original dataset
transcript = transcript.merge(average_frequency, on = 'person')

In [51]: #counts is engineered feature which will give us idea about how many interactions are ha
#groupby person , event and apply count of time and unstack with suffix
counts = transcript.groupby(by=['person', 'event'])['time'].count().unstack().add_suffi
transcript = transcript.merge(counts, on = 'person')

In [52]: #average_amount is engineered feature which will give us idea about average amount spen
#groupby person , amount and apply mean of amount and unstack with suffix
average_amount = transcript.groupby(by=['person'])['amount'].mean().to_frame().rename(c
transcript = transcript.merge(average_amount, on = 'person')

In [53]: #count_rewards is engineered feature which will give us idea about count of reward reci
#groupby person , reward and apply mean of mean and unstack with suffix
count_rewards = transcript.groupby(by=['person'])['reward'].count().to_frame().rename(c
transcript = transcript.merge(count_rewards, on = 'person')

In [54]: #total_amount is engineered feature which will give us idea about total amount spent du
#groupby person , amount and apply sum of amount and unstack with suffix
total_amount = transcript.groupby(by=['person'])['amount'].sum().to_frame().rename(colu
transcript = transcript.merge(total_amount, on = 'person')

In [55]: transcript.rename(columns={'person': 'consumer_id'}, inplace=True)

In [56]: transcript.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 306534 entries, 0 to 306533
Data columns (total 28 columns):
event                306534 non-null object
consumer_id          306534 non-null object

```

```

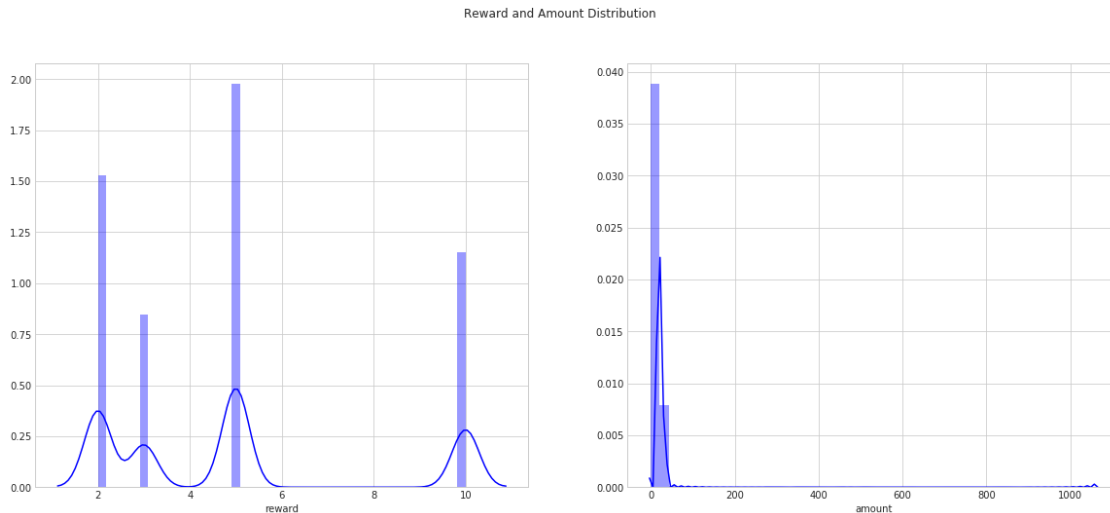
time                306534 non-null float64
period_Afternoon    306534 non-null uint8
period_Evening       306534 non-null uint8
period_Morning       306534 non-null uint8
period_Night         306534 non-null uint8
day_Friday           306534 non-null uint8
day_Monday           306534 non-null uint8
day_Saturday         306534 non-null uint8
day_Sunday           306534 non-null uint8
day_Thursday         306534 non-null uint8
day_Tuesday          306534 non-null uint8
day_Wednesday        306534 non-null uint8
offer_id             167581 non-null object
reward              33579 non-null float64
amount              138953 non-null float64
offer_completed_frequency 205346 non-null float64
offer_received_frequency 305948 non-null float64
offer_viewed_frequency 294115 non-null float64
transaction_frequency 299454 non-null float64
offer_completed_counts 254719 non-null float64
offer_received_counts 306514 non-null float64
offer_viewed_counts  305264 non-null float64
transaction_counts   303161 non-null float64
average_amount       303161 non-null float64
count_reward         306534 non-null int64
total_amount         306534 non-null float64
dtypes: float64(13), int64(1), object(3), uint8(11)
memory usage: 45.3+ MB

```

```

In [57]: #plot reward and amount distribution
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(20, 8))
sns.distplot(a=transcript['reward'][transcript['reward'].notnull()], color='b', ax=axes[0])
sns.distplot(a=transcript['amount'][transcript['amount'].notnull()], color='b', ax=axes[1])
fig.suptitle('Reward and Amount Distribution');

```



```
In [58]: transcript_later = transcript.copy()
```

```
In [59]: #drop reward and amount columns as relevant details are already captured
transcript.drop(columns=['reward', 'amount'], inplace=True)
```

```
In [60]: transcript.fillna(transcript.median(), inplace=True)
```

Merging all the dataset into a single master df

```
In [61]: df = transcript.merge(portfolio, how='left', on='offer_id' )
df = df.merge(profile, how='left', on='consumer_id')
```

```
In [62]: df.head()
```

```
Out[62]:
```

	event	consumer_id	time	period_Afternoon	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0.00	0	
1	offer viewed	78afa995795e4d85b5d9ceeca43f5fef	0.25	1	
2	transaction	78afa995795e4d85b5d9ceeca43f5fef	5.50	0	
3	offer completed	78afa995795e4d85b5d9ceeca43f5fef	5.50	0	
4	transaction	78afa995795e4d85b5d9ceeca43f5fef	6.00	0	

	period_Evening	period_Morning	period_Night	day_Friday	day_Monday	\
0	0	1	0	0	1	
1	0	0	0	0	1	
2	1	0	0	0	0	
3	1	0	0	0	0	
4	0	1	0	0	0	

	day_Saturday	day_Sunday	day_Thursday	day_Tuesday	day_Wednesday	\
0	0	0	0	0	0	
1	0	0	0	0	0	

2	1	0	0	0	0
3	1	0	0	0	0
4	0	1	0	0	0

	offer_id	offer completed_frequency \
0	9b98b8c7a33c4b65b9aebfe6a799e6d9	7.875
1	9b98b8c7a33c4b65b9aebfe6a799e6d9	7.875
2	None	7.875
3	9b98b8c7a33c4b65b9aebfe6a799e6d9	7.875
4	None	7.875

	offer received_frequency	offer viewed_frequency	transaction_frequency \
0	7.0	8.0	2.791667
1	7.0	8.0	2.791667
2	7.0	8.0	2.791667
3	7.0	8.0	2.791667
4	7.0	8.0	2.791667

	offer completed_counts	offer received_counts	offer viewed_counts \
0	3.0	4.0	4.0
1	3.0	4.0	4.0
2	3.0	4.0	4.0
3	3.0	4.0	4.0
4	3.0	4.0	4.0

	transaction_counts	average_amount	count_reward	total_amount	difficulty \
0	7.0	22.752857	3	159.27	5.0
1	7.0	22.752857	3	159.27	5.0
2	7.0	22.752857	3	159.27	NaN
3	7.0	22.752857	3	159.27	5.0
4	7.0	22.752857	3	159.27	NaN

	duration	reward	channel_email	channel_mobile	channel_social \
0	7.0	5.0	1.0	1.0	0.0
1	7.0	5.0	1.0	1.0	0.0
2	NaN	NaN	NaN	NaN	NaN
3	7.0	5.0	1.0	1.0	0.0
4	NaN	NaN	NaN	NaN	NaN

	channel_web	rew_by_diff	offer_type_bogo	offer_type_discount \
0	1.0	1.0	1.0	0.0
1	1.0	1.0	1.0	0.0
2	NaN	NaN	NaN	NaN
3	1.0	1.0	1.0	0.0
4	NaN	NaN	NaN	NaN

	offer_type_informational	age	became_member_on	income	gender_F \
0	0.0	75.0	0.756994	100000.0	1

1	0.0	75.0	0.756994	100000.0	1
2	NaN	75.0	0.756994	100000.0	1
3	0.0	75.0	0.756994	100000.0	1
4	NaN	75.0	0.756994	100000.0	1

	gender_M	gender_O	Age_group_Adult	Age_group_Old	Age_group_Young_Adult	\
0	0	0	0	1		0
1	0	0	0	1		0
2	0	0	0	1		0
3	0	0	0	1		0
4	0	0	0	1		0

	Income_group_Lower	Income_group_Middle	Income_group_Upper
0	0	0	1
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	1

```
In [63]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 306534 entries, 0 to 306533
Data columns (total 49 columns):
event                306534 non-null object
consumer_id          306534 non-null object
time                 306534 non-null float64
period_Afternoon     306534 non-null uint8
period_Evening       306534 non-null uint8
period_Morning       306534 non-null uint8
period_Night         306534 non-null uint8
day_Friday            306534 non-null uint8
day_Monday            306534 non-null uint8
day_Saturday          306534 non-null uint8
day_Sunday            306534 non-null uint8
day_Thursday          306534 non-null uint8
day_Tuesday           306534 non-null uint8
day_Wednesday        306534 non-null uint8
offer_id              167581 non-null object
offer_completed_frequency 306534 non-null float64
offer_received_frequency 306534 non-null float64
offer_viewed_frequency  306534 non-null float64
transaction_frequency  306534 non-null float64
offer_completed_counts  306534 non-null float64
offer_received_counts   306534 non-null float64
offer_viewed_counts     306534 non-null float64
transaction_counts      306534 non-null float64
average_amount         306534 non-null float64
```

```

count_reward          306534 non-null int64
total_amount          306534 non-null float64
difficulty             167581 non-null float64
duration              167581 non-null float64
reward                167581 non-null float64
channel_email         167581 non-null float64
channel_mobile        167581 non-null float64
channel_social        167581 non-null float64
channel_web           167581 non-null float64
rew_by_diff           167581 non-null float64
offer_type_bogo       167581 non-null float64
offer_type_discount   167581 non-null float64
offer_type_informational 167581 non-null float64
age                   306534 non-null float64
became_member_on      306534 non-null float64
income                306534 non-null float64
gender_F              306534 non-null uint8
gender_M              306534 non-null uint8
gender_0              306534 non-null uint8
Age_group_Adult       306534 non-null uint8
Age_group_Old         306534 non-null uint8
Age_group_Young_Adult 306534 non-null uint8
Income_group_Lower    306534 non-null uint8
Income_group_Middle   306534 non-null uint8
Income_group_Upper    306534 non-null uint8
dtypes: float64(25), int64(1), object(3), uint8(20)
memory usage: 76.0+ MB

```

Now after merging. For analysis and plotting offer_id contains values which will difficult to interpret so we change into offer{Number}

```

In [64]: #list of unique offer_id
         unique_ids = list(df['offer_id'].unique())
         #saving unique string to each offer_id
         for i in range(len(unique_ids)):
             df['offer_id'] = df['offer_id'].apply(lambda x: f'offer{i+1}' if x == unique_ids[i])

In [65]: df_analysis_1 = df.copy()

```

As offer_id is categorical variable for future model prediction it is being encoded

```

In [66]: #one-hot encoding
         offer_ids = pd.get_dummies(df['offer_id'], prefix='offer_id', prefix_sep='_')
         #concatenate with original dataset
         df = pd.concat([df, offer_ids], axis=1)
         #drop offer_id column
         df.drop(columns=['offer_id'], inplace=True)

```


6 More Data Analysis

7 Which channel is effective in acheiveing offer views?

```
In [67]: #collect offer viewed data; offer received data along with channels from df
df_analysis_1 = df[df['event'] == 'offer viewed'].iloc[:, 28:32].sum().to_frame().reset_index()
df_analysis_2 = df[df['event'] == 'offer received'].iloc[:, 28:32].sum().to_frame().reset_index()
df_analysis_1.rename(columns = {0:'count_viewed', 'index': 'channel'}, inplace=True)
df_analysis_2.rename(columns = {0:'count_received', 'index': 'channel'}, inplace=True)
#merge both dataset
df_analysis_3 = pd.merge(df_analysis_1, df_analysis_2, on = 'channel')
df_analysis_3['Percent_Viewed'] = (df_analysis_3['count_viewed']/df_analysis_3['count_received'])*100
```

```
In [68]: df_analysis_3
```

```
Out[68]:
```

	channel	count_viewed	count_received	Percent_Viewed
0	channel_email	57725.0	76277.0	75.678121
1	channel_mobile	55062.0	68609.0	80.254777
2	channel_social	42629.0	45683.0	93.314800
3	channel_web	44322.0	61001.0	72.657825

So it is clear that Social Media is most effective in acheiving offer views

8 What are the best days for offer completions?

```
In [69]: df_analysis_4
```

```
Out[69]:
```

	event	person	time	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0	
1	offer received	a03223e636434f42ac4c3df47e8bac43	0	
2	offer received	e2127556f4f64592b11af22de27a7932	0	
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0	
4	offer received	68617ca6246f4fbc85e91a2a49552598	0	
5	offer received	389bc3fa690240e798340f5a15918d5c	0	
6	offer received	c4863c7985cf408faee930f111475da3	0	
7	offer received	2eeac8d8feae4a8cad5a6af0499a211d	0	
8	offer received	aa4862eba776480b8bb9c68455b8c2e1	0	
9	offer received	31dda685af34476cad5bc968bdb01c53	0	
10	offer received	744d603ef08c4f33af5a61c8c7628d1c	0	
11	offer received	3d02345581554e81b7b289ab5e288078	0	
12	offer received	4b0da7e80e5945209a1fdddfe813dbe0	0	
13	offer received	c27e0d6ab72c455a8bb66d980963de60	0	
14	offer received	d53717f5400c4e84affdaeda9dd926b3	0	
15	offer received	f806632c011441378d4646567f357a21	0	
16	offer received	d058f73bf8674a26a95227db098147b1	0	
17	offer received	65aba5c617294649aeb624da249e1ee5	0	
18	offer received	ebe7ef46ea6f4963a7dd49f501b26779	0	
19	offer received	1e9420836d554513ab90eba98552d0a9	0	

20	offer received	868317b9be554cb18e50bc68484749a2	0
21	offer received	f082d80f0aac47a99173ba8ef8fc1909	0
22	offer received	102e9454054946fda62242d2e176fdce	0
23	offer received	4beeb3ed64dd4898b0edf2f6b67426d3	0
24	offer received	9f30b375d7bd4c62a884ffe7034e09ee	0
25	offer received	25c906289d154b66bf579693f89481c9	0
26	offer received	6e014185620b49bd98749f728747572f	0
27	offer received	02c083884c7d45b39cc68e1314fec56c	0
28	offer received	c0d210398dee4a0895b24444a5fcd1d2	0
29	offer received	8be4463721e14d7fa600686bf8c8b2ed	0
...
306504	transaction	8524d450673b4c24869b6c94380006de	714
306505	transaction	b895c57e8cd047a8872ce02aa54759d6	714
306506	offer completed	b895c57e8cd047a8872ce02aa54759d6	714
306507	offer viewed	8dda575c2a1d44b9ac8e8b07b93d1f8e	714
306508	transaction	8431c16f8e1d440880db371a68f82dd0	714
306509	offer completed	8431c16f8e1d440880db371a68f82dd0	714
306510	transaction	ba620885e51c4b0ea64a4f61daad494f	714
306511	transaction	a1a8f40407c444cc848468275308958a	714
306512	transaction	8d80970192fa496f99d6b45c470a4b60	714
306513	transaction	bde275066f3c4fa0bff3093e3b866a2c	714
306514	transaction	f1e4fd36e5a0446f83861308bddf6945	714
306515	transaction	0b64be3b241c4407a5c9a71781173829	714
306516	transaction	86d03d35d7e0434b935e7743e83be3a0	714
306517	transaction	3408fd05c781401f8442fb6dba9c7	714
306518	transaction	1593d617fac246ef8e50dbb0ffd77f5f	714
306519	transaction	f1b31d07b5d84f69a2d5f1d07843989e	714
306520	transaction	2ce987015ec0404a97ba333e8e814090	714
306521	transaction	2e33545f0a764d27b2ccff95fc8d72c4	714
306522	transaction	d1c4500ace2e45e9a45d3cd2fccac8d8	714
306523	transaction	b65affd9e07346a1906364a396950e3d	714
306524	transaction	d613ca9c59dd42f497bdbf6178da54a7	714
306525	transaction	eec70ab28af74a22a4aeb889c0317944	714
306526	transaction	24f56b5e1849462093931b164eb803b5	714
306527	offer completed	24f56b5e1849462093931b164eb803b5	714
306528	transaction	5ca2620962114246ab218fc648eb3934	714
306529	transaction	b3a1272bc9904337b331bf348c3e8c17	714
306530	transaction	68213b08d99a4ae1b0dcb72aebd9aa35	714
306531	transaction	a00058cf10334a308c68e7631c529907	714
306532	transaction	76ddbd6576844afe811f1a3c0fbb5bec	714
306533	transaction	c02b10e8752c4d8e9b73f918558531f7	714

		value	period_Afternoon	\
0	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}		0	
1	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}		0	
2	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}		0	
3	{'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}		0	
4	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}		0	

5	{'offer id': 'f19421c1d4aa40978ebb69ca19b0e20d'}	0
6	{'offer id': '2298d6c36e964ae4a3e7e9706d1fb8c2'}	0
7	{'offer id': '3f207df678b143eea3cee63160fa8bed'}	0
8	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
9	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
10	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
11	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
12	{'offer id': 'ae264e3637204a6fb9bb56bc8210ddfd'}	0
13	{'offer id': '3f207df678b143eea3cee63160fa8bed'}	0
14	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
15	{'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}	0
16	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
17	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0
18	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	0
19	{'offer id': 'ae264e3637204a6fb9bb56bc8210ddfd'}	0
20	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0
21	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	0
22	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}	0
23	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0
24	{'offer id': '2298d6c36e964ae4a3e7e9706d1fb8c2'}	0
25	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0
26	{'offer id': 'f19421c1d4aa40978ebb69ca19b0e20d'}	0
27	{'offer id': 'ae264e3637204a6fb9bb56bc8210ddfd'}	0
28	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	0
29	{'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}	0
...
306504	{'amount': 4.89}	0
306505	{'amount': 4.48}	0
306506	{'offer_id': 'fafdcd668e3743c1bb461111dcafc2a4...}	0
306507	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
306508	{'amount': 1.19}	0
306509	{'offer_id': 'fafdcd668e3743c1bb461111dcafc2a4...}	0
306510	{'amount': 14.31}	0
306511	{'amount': 2.37}	0
306512	{'amount': 6.92}	0
306513	{'amount': 12.73}	0
306514	{'amount': 8.2}	0
306515	{'amount': 2.6}	0
306516	{'amount': 9.2}	0
306517	{'amount': 11.7}	0
306518	{'amount': 40.67}	0
306519	{'amount': 31.13}	0
306520	{'amount': 1.6400000000000001}	0
306521	{'amount': 17.35}	0
306522	{'amount': 4.42}	0
306523	{'amount': 18.35}	0
306524	{'amount': 25.14}	0
306525	{'amount': 43.58}	0

306526		{'amount': 22.64}	0
306527	{'offer_id': 'fafdcd668e3743c1bb461111dcafc2a4...		0
306528		{'amount': 2.2}	0
306529		{'amount': 1.5899999999999999}	0
306530		{'amount': 9.53}	0
306531		{'amount': 3.61}	0
306532		{'amount': 3.5300000000000002}	0
306533		{'amount': 4.05}	0

	period_Evening	period_Morning	period_Night	day	day_Friday	\
0	0	1	0	Monday	0	
1	0	1	0	Monday	0	
2	0	1	0	Monday	0	
3	0	1	0	Monday	0	
4	0	1	0	Monday	0	
5	0	1	0	Monday	0	
6	0	1	0	Monday	0	
7	0	1	0	Monday	0	
8	0	1	0	Monday	0	
9	0	1	0	Monday	0	
10	0	1	0	Monday	0	
11	0	1	0	Monday	0	
12	0	1	0	Monday	0	
13	0	1	0	Monday	0	
14	0	1	0	Monday	0	
15	0	1	0	Monday	0	
16	0	1	0	Monday	0	
17	0	1	0	Monday	0	
18	0	1	0	Monday	0	
19	0	1	0	Monday	0	
20	0	1	0	Monday	0	
21	0	1	0	Monday	0	
22	0	1	0	Monday	0	
23	0	1	0	Monday	0	
24	0	1	0	Monday	0	
25	0	1	0	Monday	0	
26	0	1	0	Monday	0	
27	0	1	0	Monday	0	
28	0	1	0	Monday	0	
29	0	1	0	Monday	0	
...
306504	0	0	1	Tuesday	0	
306505	0	0	1	Tuesday	0	
306506	0	0	1	Tuesday	0	
306507	0	0	1	Tuesday	0	
306508	0	0	1	Tuesday	0	
306509	0	0	1	Tuesday	0	
306510	0	0	1	Tuesday	0	

306511	0	0	1	Tuesday	0
306512	0	0	1	Tuesday	0
306513	0	0	1	Tuesday	0
306514	0	0	1	Tuesday	0
306515	0	0	1	Tuesday	0
306516	0	0	1	Tuesday	0
306517	0	0	1	Tuesday	0
306518	0	0	1	Tuesday	0
306519	0	0	1	Tuesday	0
306520	0	0	1	Tuesday	0
306521	0	0	1	Tuesday	0
306522	0	0	1	Tuesday	0
306523	0	0	1	Tuesday	0
306524	0	0	1	Tuesday	0
306525	0	0	1	Tuesday	0
306526	0	0	1	Tuesday	0
306527	0	0	1	Tuesday	0
306528	0	0	1	Tuesday	0
306529	0	0	1	Tuesday	0
306530	0	0	1	Tuesday	0
306531	0	0	1	Tuesday	0
306532	0	0	1	Tuesday	0
306533	0	0	1	Tuesday	0

	day_Monday	day_Saturday	day_Sunday	day_Thrusday	day_Tuesday	\
0	1	0	0	0	0	
1	1	0	0	0	0	
2	1	0	0	0	0	
3	1	0	0	0	0	
4	1	0	0	0	0	
5	1	0	0	0	0	
6	1	0	0	0	0	
7	1	0	0	0	0	
8	1	0	0	0	0	
9	1	0	0	0	0	
10	1	0	0	0	0	
11	1	0	0	0	0	
12	1	0	0	0	0	
13	1	0	0	0	0	
14	1	0	0	0	0	
15	1	0	0	0	0	
16	1	0	0	0	0	
17	1	0	0	0	0	
18	1	0	0	0	0	
19	1	0	0	0	0	
20	1	0	0	0	0	
21	1	0	0	0	0	
22	1	0	0	0	0	

23	1	0	0	0	0
24	1	0	0	0	0
25	1	0	0	0	0
26	1	0	0	0	0
27	1	0	0	0	0
28	1	0	0	0	0
29	1	0	0	0	0
...
306504	0	0	0	0	1
306505	0	0	0	0	1
306506	0	0	0	0	1
306507	0	0	0	0	1
306508	0	0	0	0	1
306509	0	0	0	0	1
306510	0	0	0	0	1
306511	0	0	0	0	1
306512	0	0	0	0	1
306513	0	0	0	0	1
306514	0	0	0	0	1
306515	0	0	0	0	1
306516	0	0	0	0	1
306517	0	0	0	0	1
306518	0	0	0	0	1
306519	0	0	0	0	1
306520	0	0	0	0	1
306521	0	0	0	0	1
306522	0	0	0	0	1
306523	0	0	0	0	1
306524	0	0	0	0	1
306525	0	0	0	0	1
306526	0	0	0	0	1
306527	0	0	0	0	1
306528	0	0	0	0	1
306529	0	0	0	0	1
306530	0	0	0	0	1
306531	0	0	0	0	1
306532	0	0	0	0	1
306533	0	0	0	0	1

	day_Wednesday
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	0

8	0
9	0
10	0
11	0
12	0
13	0
14	0
15	0
16	0
17	0
18	0
19	0
20	0
21	0
22	0
23	0
24	0
25	0
26	0
27	0
28	0
29	0
...	...
306504	0
306505	0
306506	0
306507	0
306508	0
306509	0
306510	0
306511	0
306512	0
306513	0
306514	0
306515	0
306516	0
306517	0
306518	0
306519	0
306520	0
306521	0
306522	0
306523	0
306524	0
306525	0
306526	0
306527	0
306528	0

```

306529      0
306530      0
306531      0
306532      0
306533      0

```

```
[306534 rows x 16 columns]
```

```

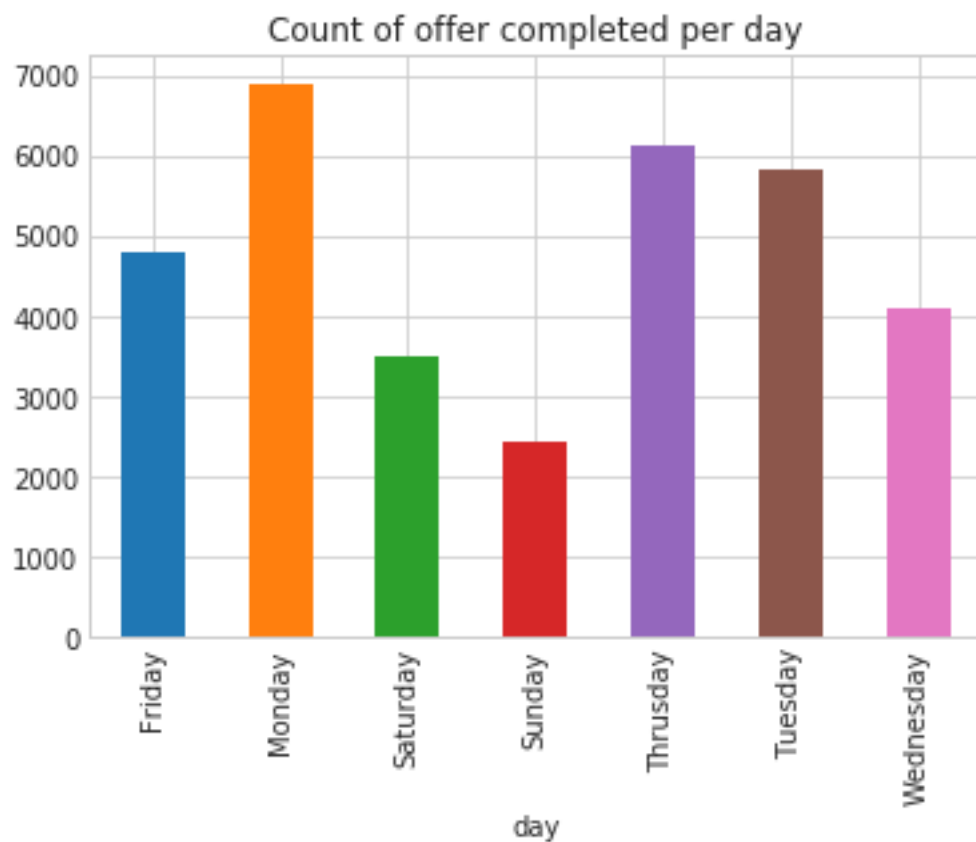
In [70]: #groupby day and count events
df_analysis_5 = df_analysis_4[df_analysis_4['event'] == 'offer completed'].groupby(by=

```

```

In [71]: df_analysis_5.plot(kind='bar');
plt.title('Count of offer completed per day');

```



It is clear that Mondays are best followed by Thrusay and Tuesday

9 In what period is most offers completed?

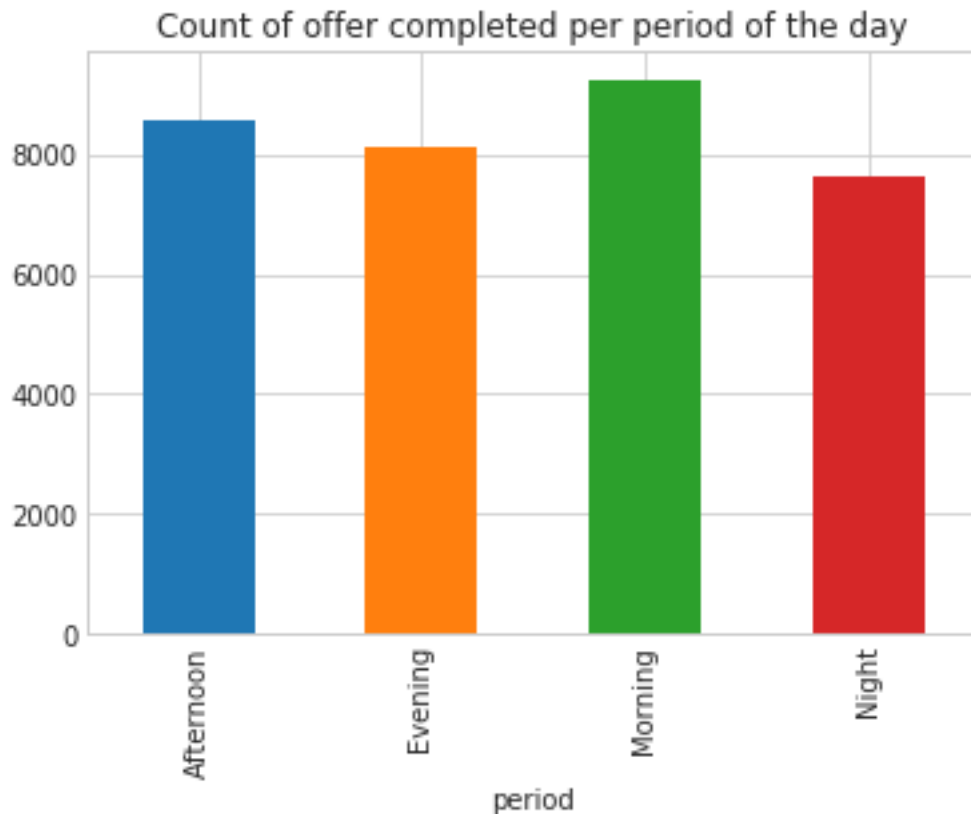
```

In [72]: #groupby period and count event
df_analysis_7 = df_analysis_6[df_analysis_6['event'] == 'offer completed'].groupby(by=

```



```
In [73]: df_analysis_7.plot(kind='bar');
plt.title('Count of offer completed per period of the day');
```



It is clear that Mornings have most offer completion

10 Predictive Model

11 What features contribute most for viewing the offer?

```
In [74]: df1 = df[((df['event']=='offer received')| (df['event']=='offer viewed'))]
```

```
In [75]: #drop the columns with no use and with directly contains information of target variable
X1 = df1.drop(columns=['event', 'consumer_id', 'offer viewed_counts', 'offer viewed_frequency'])
Y1 = df1['event'].apply(lambda x: 0 if (x == 'offer received') else 1).reset_index()['event']
```

```
In [76]: X1.shape
```

```
Out[76]: (134002, 55)
```

Below function splits the data into test and train. Gradient Boosting Classifier is used because because of its efficacy in dealing with large complex tabular datasets. Finally feature importance will be plotted and model will be returned

```

In [77]: def feature_importance(X, Y):
    '''
    Plots important features and return the machine learning model
    Input:
    X: Dataframe containing the input columns
    Y: Series containing the output column
    Output:
    Returns model fitted on Input data
    '''
    #Split the data
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_s
    #initiate the model
    model = GradientBoostingClassifier()
    #Cross Validation
    kfold = KFold(n_splits=5)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold, scoring='roc_auc',
    print("Mean: %f ;standard deviation: %f" % (cv_results.mean(), cv_results.std()))
    #fit the data
    model.fit(X_train, Y_train)
    #predict the the test data
    predictions = model.predict(X_test)
    print("Accuracy: %f" % accuracy_score(predictions ,Y_test))
    print("AUC : %f" % roc_auc_score(predictions ,Y_test))
    #store feature importance
    feature_imp = pd.DataFrame(sorted(zip(model.feature_importances_,X_train.columns)),
    plt.figure(figsize=(25, 10))
    sns.barplot(x="Value", y="Feature", data=feature_imp.sort_values(by="Value", ascend
    plt.tight_layout()
    plt.show()
    return model

```

```

In [78]: model1 = feature_importance(X1, Y1)

```

```

Mean: 0.923400 ;standard deviation: 0.001714

```

```


Accuracy: 0.903586

```

```

AUC : 0.927660

```



Above Visualization have huge benefit. Following are the insights that can be derived: - Morning period play important role for offer viewing - Social Media advertisement is highly effective in offer viewing - Income and Average amount spend play role in viewing. - Next the day of the week has huge impact on whether the offer is viewed and left alone. This will help in targeting most important offers on particular days

12 What features contribute most for completing the offer?

```
In [79]: df2 = df[((df['event']=='offer viewed')| (df['event']=='offer completed')) | (df['event']
```

```
In [80]: df2.head()
```

```
Out[80]:
```

	event	consumer_id	time	period_Afternoon	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0.00	0	
1	offer viewed	78afa995795e4d85b5d9ceeca43f5fef	0.25	1	
3	offer completed	78afa995795e4d85b5d9ceeca43f5fef	5.50	0	
5	offer received	78afa995795e4d85b5d9ceeca43f5fef	7.00	0	
6	offer viewed	78afa995795e4d85b5d9ceeca43f5fef	9.00	0	

	period_Evening	period_Morning	period_Night	day_Friday	day_Monday	\
0	0	1	0	0	1	
1	0	0	0	0	1	
3	1	0	0	0	0	
5	0	1	0	0	1	
6	0	1	0	0	0	

	day_Saturday	day_Sunday	day_Thursday	day_Tuesday	day_Wednesday	\
0	0	0	0	0	0	
1	0	0	0	0	0	
3	1	0	0	0	0	
5	0	0	0	0	0	
6	0	0	0	0	1	

	offer_completed_frequency	offer_received_frequency	\
0	7.875	7.0	
1	7.875	7.0	
3	7.875	7.0	
5	7.875	7.0	
6	7.875	7.0	

	offer_viewed_frequency	transaction_frequency	offer_completed_counts	\
0	8.0	2.791667	3.0	
1	8.0	2.791667	3.0	
3	8.0	2.791667	3.0	
5	8.0	2.791667	3.0	
6	8.0	2.791667	3.0	

	offer_received_counts	offer_viewed_counts	transaction_counts	\
0	4.0	4.0	7.0	
1	4.0	4.0	7.0	
3	4.0	4.0	7.0	
5	4.0	4.0	7.0	
6	4.0	4.0	7.0	

	average_amount	count_reward	total_amount	difficulty	duration	reward	\
0	22.752857	3	159.27	5.0	7.0	5.0	
1	22.752857	3	159.27	5.0	7.0	5.0	
3	22.752857	3	159.27	5.0	7.0	5.0	
5	22.752857	3	159.27	0.0	3.0	0.0	
6	22.752857	3	159.27	0.0	3.0	0.0	

	channel_email	channel_mobile	channel_social	channel_web	rew_by_diff	\
0	1.0	1.0	0.0	1.0	1.0	
1	1.0	1.0	0.0	1.0	1.0	
3	1.0	1.0	0.0	1.0	1.0	
5	1.0	1.0	1.0	0.0	0.0	
6	1.0	1.0	1.0	0.0	0.0	

	offer_type_bogo	offer_type_discount	offer_type_informational	age	\
0	1.0	0.0	0.0	75.0	
1	1.0	0.0	0.0	75.0	
3	1.0	0.0	0.0	75.0	
5	0.0	0.0	1.0	75.0	
6	0.0	0.0	1.0	75.0	

	became_member_on	income	gender_F	gender_M	gender_0	Age_group_Adult	\
0	0.756994	100000.0	1	0	0	0	
1	0.756994	100000.0	1	0	0	0	
3	0.756994	100000.0	1	0	0	0	
5	0.756994	100000.0	1	0	0	0	

6	0.756994	100000.0	1	0	0	0
---	----------	----------	---	---	---	---

	Age_group_Old	Age_group_Young_Adult	Income_group_Lower	\
0	1	0	0	
1	1	0	0	
3	1	0	0	
5	1	0	0	
6	1	0	0	

	Income_group_Middle	Income_group_Upper	offer_id_offer1	offer_id_offer10	\
0	0	1	1	0	
1	0	1	1	0	
3	0	1	1	0	
5	0	1	0	0	
6	0	1	0	0	

	offer_id_offer11	offer_id_offer2	offer_id_offer3	offer_id_offer4	\
0	0	0	0	0	
1	0	0	0	0	
3	0	0	0	0	
5	0	0	1	0	
6	0	0	1	0	

	offer_id_offer5	offer_id_offer6	offer_id_offer7	offer_id_offer8	\
0	0	0	0	0	
1	0	0	0	0	
3	0	0	0	0	
5	0	0	0	0	
6	0	0	0	0	

	offer_id_offer9
0	0
1	0
3	0
5	0
6	0


```
In [81]: #drop the columns with no use and with directly contains information of target variable
X2 = df2.drop(columns=['event', 'consumer_id'])
Y2 = df2['event'].apply(lambda x: 1 if (x == 'offer completed') else 0).reset_index()['event']
```

```
In [82]: model2 = feature_importance(X2, Y2)
```

Mean: 0.912877 ;standard deviation: 0.001245

Accuracy: 0.863055

AUC : 0.790204



Following insights can be derived from above plot: - Just like in the case of offer viewing, Morning period play important role in offer completion - Next, Looking at count reward and offer_completed_count indicates that people keep buying as they receive rewards and become loyal customers - It is surprising to offer viewed count and offer viewed frequency indicating that personalized frequent offers sent will impact offer completion - Thursday is also important day for offer completion

13 Evaluation

Results are showcased in the medium blog post link: <https://anudeeppeela9.medium.com/starbucks-capstone-challenge-e901baeff5d2>

13.0.1 Refinement Documentation

```
In [85]: #Hyperparameter tuning
def tuning(X, Y):
    """
    Prints the score after tuning and return the tuned model
    Input:
    X: Dataframe containing the input columns
    Y: Series containing the output column
    Output:
    Returns model fitted on Input data
    """
    #Split the data
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_s
    #Hyperparameters
    num_estimators = [250, 500]
    learn_rates = [0.05, 0.1]
    #initiate the model
    model = GradientBoostingClassifier()
    #parameter grid
```

```

param_grid = {'n_estimators': num_estimators, 'learning_rate': learn_rates}
kfold = KFold(n_splits=3)
#gridsearchcv
grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring='roc_auc', cv=kfold)
grid_result = grid.fit(X_train, Y_train)
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
#predict the the test data
predictions = grid.predict(X_test)
print("Accuracy: %f" % accuracy_score(predictions, Y_test))
print("AUC : %f" % roc_auc_score(predictions, Y_test))
return model

```

In [86]: tuning(X1, Y1)

```

Best: 0.923374 using {'learning_rate': 0.05, 'n_estimators': 250}
Accuracy: 0.903884
AUC : 0.927851

```

```

Out[86]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=3,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)

```

In [87]: tuning(X2, Y2)

```

Best: 0.917053 using {'learning_rate': 0.1, 'n_estimators': 500}
Accuracy: 0.870364
AUC : 0.797834

```

```

Out[87]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=3,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)

```

In [88]: *#logistic regression*

```

from sklearn.linear_model import LogisticRegression
def feature_importance_1(X, Y):
    """
    Prints accuray and outputs model

```

```

Input:
X: Dataframe containing the input columns
Y: Series containing the output column
Output:
Returns model fitted on Input data
'''
#Split the data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_s
#initiate the model
model = LogisticRegression()
#fit the data
model.fit(X_train, Y_train)
#predict the the test data
predictions = model.predict(X_test)
print("Accuracy: %f" % accuracy_score(predictions ,Y_test))
print("AUC : %f" % roc_auc_score(predictions ,Y_test))
return model

```

In [89]: feature_importance_1(X1, Y1)

Accuracy: 0.872094
AUC : 0.908228

```

Out[89]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)

```

In [90]: feature_importance_1(X2, Y2)

Accuracy: 0.797446
AUC : 0.580042

```

Out[90]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)

```

In [92]: !pip install lightgbm

Collecting lightgbm

Downloading <https://files.pythonhosted.org/packages/18/b2/fff8370f48549ce223f929fe8cab4ee6bf28>
100% || 2.0MB 8.3MB/s eta 0:00:01

Requirement already satisfied: wheel in /opt/conda/lib/python3.6/site-packages (from lightgbm) (
Requirement already satisfied: numpy in /opt/conda/lib/python3.6/site-packages (from lightgbm) (
Requirement already satisfied: scikit-learn!=0.22.0 in /opt/conda/lib/python3.6/site-packages (f
Requirement already satisfied: scipy in /opt/conda/lib/python3.6/site-packages (from lightgbm) (

Installing collected packages: lightgbm
Successfully installed lightgbm-3.2.1

```
In [93]: #light lgb
import lightgbm as lgb
def feature_importance_2(X, Y):
    '''
    Returns model fitted on Input data
    Input:
    X: Dataframe containing the input columns
    Y: Series containing the output column
    Output:
    Returns model fitted on Input data
    '''
    #Split the data
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_s
    #initiate the model
    model = lgb.LGBMClassifier()
    #fit the data
    model.fit(X_train, Y_train)
    #predict the the test data
    predictions = model.predict(X_test)
    print("Accuracy: %f" % accuracy_score(predictions ,Y_test))
    print("AUC : %f" % roc_auc_score(predictions ,Y_test))
    return model
```

```
In [94]: feature_importance_2(X1, Y1)
```

Accuracy: 0.903847
AUC : 0.927791

```
Out[94]: LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
    importance_type='split', learning_rate=0.1, max_depth=-1,
    min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,
    n_estimators=100, n_jobs=-1, num_leaves=31, objective=None,
    random_state=None, reg_alpha=0.0, reg_lambda=0.0, silent=True,
    subsample=1.0, subsample_for_bin=200000, subsample_freq=0)
```

```
In [95]: feature_importance_2(X2, Y2)
```

Accuracy: 0.870066
AUC : 0.796191

```
Out[95]: LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
    importance_type='split', learning_rate=0.1, max_depth=-1,
    min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,
```

```
n_estimators=100, n_jobs=-1, num_leaves=31, objective=None,  
random_state=None, reg_alpha=0.0, reg_lambda=0.0, silent=True,  
subsample=1.0, subsample_for_bin=200000, subsample_freq=0)
```

```
In [ ]:
```